

IST 597: Foundations of Deep Learning

Assignment 10

Instructor : Dr. C. Lee. Giles

TA : Neisarg Dave

Due Date : Tuesday, September 12, 2023

Course Policy

- Carefully read all the instructions before you start working on the assignment
- Give maximum explanation for each sub-problem. Please avoid single-line answers; submissions without any explanations will receive 0 points.
- Assignments are due before class at 02:29 pm. Please check the due date on Canvas.
- Late exercises will receive 50% credits for the first 24 hours and no credits thereafter.
- All exercise solutions must be turned in, even if late. Failure to do so can result in a deferred grade.
- All source materials must be cited. The University Academic Code of Conduct will be strictly enforced.
- All queries related to Assignment should have a subject line *IST597 : Assignment_01 Queries*

Assignment Instructions:

- The submission for this assignment must be a zipped folder: `{name}_assignment_10.zip`
- The folder must contain two files:
 1. `{name}_assignment_10.pdf`: Answers for all theory questions and explanations for all problems.
 2. `{name}_assignment_10.py`: Codes for programming problems

A template Python file is provided. Feel free to make any suitable changes.

1 Linear Layer

A linear layer in a neural network is a affine transformation of its inputs, defined as :

$$z = Wx + b$$

where x and z are inputs and output respectively and W and b are parameters. W is called weight matrix and b is bias.

2 Multilayer Perceptron

A multilayer perceptron is constructed by placing multiple linear layers, combined with an activation function. Activation function provides non-linearity to the multilayer perceptron. A good example of activation function is Rectified Linear Unit (ReLU) function defined as :

$$ReLU(x) = \begin{cases} x, & x \geq 0 \\ 0, & x < 0 \end{cases}$$

An MLP of L layers can be written as:

$$\begin{aligned}\hat{z}^{(l)} &= W^{(l)} z^{(l-1)} + b \\ z^{(l)} &= Activation(\hat{z}^{(l)}) \\ z^{(0)} &= x, \quad l \in [0, L-1]\end{aligned}$$

3 Loss Function : Cross Entropy

Cross Entropy Loss for a multiclass classification problem is defined as:

$$\mathcal{L}(W) = ce_loss = \frac{1}{\#samples} \sum_i^{\#samples} \sum_j^{\#classes} -y_{ij} \cdot \log(p_{ij})$$

where,

$$p_{ij} = \frac{e^{h_{ij}}}{\sum_k e^{h_{ik}}} \quad \text{and} \quad h_i = z_i^{(L-1)}$$

4 Gradient Calculation for Backpropagation

For a given affine transformation:

$$z^{(l)} = W^{(l)} x^{(l-1)} + b$$

We can calculate all the partial derivations in the following fashion

$$\frac{\partial z^{(l)}}{\partial x^{(l-1)}} = W^{(l)}$$

$$\frac{\partial z^{(l)}}{\partial W^{(l)}} = x^{(l-1)}$$

$$\frac{\partial z^{(l)}}{\partial b^{(l)}} = 1$$

To write the **backward** function, we need one more step. **backward** function takes **grad_output** as its input. It is the backward flow of gradients coming from the final output \hat{y} and can be mathematically expressed as:

$$grad_output = \frac{\partial \hat{y}}{\partial z^{(l)}}$$

backward functions gives out three values:

$$grad_x = \frac{\partial \hat{y}}{\partial x^{(l-1)}} = \frac{\partial \hat{y}}{\partial z^{(l)}} \frac{\partial z^{(l)}}{\partial x^{(l-1)}}$$

$$grad_w = \frac{\partial \hat{y}}{\partial W^{(l)}} = \frac{\partial \hat{y}}{\partial z^{(l)}} \frac{\partial z^{(l)}}{\partial W^{(l)}}$$

$$grad.b = \frac{\partial \hat{y}}{\partial b^{(l)}} = \frac{\partial \hat{y}}{\partial z^{(l)}} \frac{\partial z^{(l)}}{\partial b^{(l)}}$$

In the given code template update the **forward** and backward functions of class **LinearFunction** to create an affine transformation layer.

5 Parameter Update Equation

Backpropagation Algorithm is used for calculating the gradients of loss function with respect to the parameters ($\nabla_W \mathcal{L}$). The parameters then can be updated at the end of each iteration:

$$\hat{W} = W - \alpha \nabla_W \mathcal{L}$$

6 Assignment

In this assignment, you are required to do the following:

1. Create an MLP 3 marks
 - (a) Write the **forward** function for affine transformation
 - (b) Write the **backward** function (to calculate gradients) for the affine transformation
 - (c) Create a Neural Network composed of multiple affine transformations.
2. Train the Neural Network on Fashion MNIST Dataset 4 marks
 - (a) Write code for Loss Function.
 - (b) Write code for the parameter update routine
 - (c) Write code to calculate loss on validation set. Validation should be run after every epoch.
 - (d) Plot Train and Validation Loss Curves using Matplotlib
3. Testing the Neural Network 2 marks
 - (a) Change the **seed** and train the neural network 5 times
 - (b) Report the mean and variance over 5 trials on the following metrics on test set:
 - i. Accuracy
 - ii. Precision
 - iii. Recall
 - iv. F1
 - v. ROC AUC
4. Change the number of parameters in the Neural Network 3 marks
 - (a) Train an underfit model
 - (b) Train a overfit model
 - (c) Demonstrate the difference between overfitting and underfitting